

Images reading and processing

```
In [1]: import cv2
import os
import matplotlib.pyplot as plt
import numpy as np
```

```
In [2]: # to read images from each file and assigning it a label
def load_images_from_folder(folder):
    images = []
    lables=[]
    c=-1
    d={}
    for file in os.listdir(folder):
        c+=1
        d[file]=c
        for filename in os.listdir(os.path.join(folder,file)):
            img = cv2.imread(os.path.join(folder,file,filename))
            if img is not None:
                images.append(img)
                lables.append([c])
    return images,lables
```

```
In [3]: folder='dataset'


images,lables=load_images_from_folder(folder)
```

```
In [4]: #checking the length of the data read
print(len(images))
print(len(lables))
```

```
2007
2007
```

```
In [5]: #checking for the images
plt.imshow(images[0])
```

```
Out[5]: <matplotlib.image.AxesImage at 0x29380572358>
```




```
In [6]: #images shape  
images[0].shape
```

```
Out[6]: (332, 183, 3)
```

```
In [7]: img = cv2.resize(images[0], (150, 250))  
  
plt.imshow(img)
```

```
Out[7]: <matplotlib.image.AxesImage at 0x29383b473c8>
```




As every image is of different size we will resize them to a feasible size and make the size of each image similar.

```
In [8]: #resizing the images  
cleaned_images=[]  
for i in range(len(images)):  
    #image = cv2.cvtColor(images[i], cv2.COLOR_RGB2GRAY)  
    #image = cv2.cvtColor(images[i], cv2.COLOR_BGR2GRAY)  
    cleaned_images.append(cv2.resize(images[i], (150, 250)))
```

```
In [9]: #splitting the data into two parts that is test and train with train having 20% of  
the total data  
from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test = train_test_split(cleaned_images, labels, test_si  
ze=0.2, random_state=42)
```


```
In [10]: #checking if every this is right upto this  
plt.imshow(X_train[0])
```

```
Out[10]: <matplotlib.image.AxesImage at 0x293ac765ac8>
```



```
In [11]: plt.imshow(X_test[0])
```

```
Out[11]: <matplotlib.image.AxesImage at 0x293b06ac358>
```



```
In [12]: print(type(X_test))  
print(type(X_train))
```

```
<class 'list'>  
<class 'list'>
```

```
In [13]: #converting them to numpy array type so that it may be further used to as type conversions  
X_test=np.array(X_test)  
X_train=np.array(X_train)  
  
X_train = X_train.astype('float32')/255  
X_test = X_test.astype('float32')/255
```

```
In [14]: print(X_train.shape)  
  
print(X_test.shape)  
  
(1605, 250, 150, 3)  
(402, 250, 150, 3)
```

so far so good

```
In [15]: from keras.utils import np_utils
import keras

# one-hot encode the labels
num_classes = len(np.unique(y_train))
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

# break training set into training and validation sets
(X_train, X_valid) = X_train[100:], X_train[:100]
(y_train, y_valid) = y_train[100:], y_train[:100]

# print shape of training set
print('x_train shape:', X_train.shape)

# print number of training, validation, and test images
print(X_train.shape[0], 'train samples')
print(X_test.shape[0], 'test samples')
print(X_valid.shape[0], 'validation samples')
```

C:\Users\Akarsh Somani\Anaconda3\lib\site-packages\h5py__init__.py:36: FutureWarning: Conversion of the second argument of issubdtype from `float` to `np.float` is deprecated. In future, it will be treated as `np.float64 == np.dtype(float).type`.

from ._conv import register_converters as _register_converters
Using TensorFlow backend.

```
x_train shape: (1505, 250, 150, 3)
1505 train samples
402 test samples
100 validation samples
```

```
In [16]: #Our CNN model
from keras.models import Sequential
from keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout

model = Sequential()
model.add(Conv2D(filters=16, kernel_size=3, padding='same', activation='relu',
                  input_shape=(250, 150, 3)))
model.add(MaxPooling2D(pool_size=2))
model.add(Conv2D(filters=32, kernel_size=2, padding='same', activation='relu'))
model.add(MaxPooling2D(pool_size=2))
model.add(Dropout(0.2))
model.add(Conv2D(filters=64, kernel_size=2, padding='same', activation='relu'))
model.add(MaxPooling2D(pool_size=2))
model.add(Dropout(0.2))
model.add(Flatten())
model.add(Dense(500, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(13, activation='softmax'))

model.summary()
```

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 250, 150, 16)	448
max_pooling2d_1 (MaxPooling2D)	(None, 125, 75, 16)	0
conv2d_2 (Conv2D)	(None, 125, 75, 32)	2080
max_pooling2d_2 (MaxPooling2D)	(None, 62, 37, 32)	0
dropout_1 (Dropout)	(None, 62, 37, 32)	0
conv2d_3 (Conv2D)	(None, 62, 37, 64)	8256
max_pooling2d_3 (MaxPooling2D)	(None, 31, 18, 64)	0
dropout_2 (Dropout)	(None, 31, 18, 64)	0
flatten_1 (Flatten)	(None, 35712)	0
dense_1 (Dense)	(None, 500)	17856500
dropout_3 (Dropout)	(None, 500)	0
dense_2 (Dense)	(None, 13)	6513

=====
 Total params: 17,873,797
 Trainable params: 17,873,797
 Non-trainable params: 0
 =====

```
In [17]: # compile the model
model.compile(loss='categorical_crossentropy', optimizer='adam',
              metrics=['accuracy'])
```

```
In [18]: from keras.callbacks import ModelCheckpoint

# train the model
checkpointer = ModelCheckpoint(filepath='model.weights.best.hdf5', verbose=1,
                               save_best_only=True)
hist = model.fit(X_train, y_train, batch_size=100, epochs=20,
                 validation_data=(X_valid, y_valid), callbacks=[checkpointer], shuffle=True
                 )
```

```
Train on 1505 samples, validate on 100 samples
Epoch 1/20
1505/1505 [=====] - 12s 8ms/step - loss: 5.3094 - acc: 0.0804 - val_loss: 2.5123 - val_acc: 0.1400

Epoch 00001: val_loss improved from inf to 2.51231, saving model to model.weights.best.hdf5
Epoch 2/20
1505/1505 [=====] - 7s 5ms/step - loss: 2.5759 - acc: 0.0817 - val_loss: 2.5599 - val_acc: 0.1700

Epoch 00002: val_loss did not improve from 2.51231
Epoch 3/20
1505/1505 [=====] - 7s 5ms/step - loss: 2.5580 - acc: 0.1110 - val_loss: 2.5603 - val_acc: 0.0800

Epoch 00003: val_loss did not improve from 2.51231
Epoch 4/20
1505/1505 [=====] - 7s 5ms/step - loss: 2.5388 - acc: 0.1236 - val_loss: 2.5308 - val_acc: 0.1500

Epoch 00004: val_loss did not improve from 2.51231
Epoch 5/20
1505/1505 [=====] - 7s 5ms/step - loss: 2.4933 - acc: 0.1548 - val_loss: 2.5004 - val_acc: 0.0900

Epoch 00005: val_loss improved from 2.51231 to 2.50039, saving model to model.weights.best.hdf5
Epoch 6/20
1505/1505 [=====] - 7s 5ms/step - loss: 2.4038 - acc: 0.2053 - val_loss: 2.4175 - val_acc: 0.2200

Epoch 00006: val_loss improved from 2.50039 to 2.41755, saving model to model.weights.best.hdf5
Epoch 7/20
1505/1505 [=====] - 7s 5ms/step - loss: 2.3000 - acc: 0.2319 - val_loss: 2.4726 - val_acc: 0.1600

Epoch 00007: val_loss did not improve from 2.41755
Epoch 8/20
1505/1505 [=====] - 7s 5ms/step - loss: 2.2102 - acc: 0.2757 - val_loss: 2.3937 - val_acc: 0.2100

Epoch 00008: val_loss improved from 2.41755 to 2.39367, saving model to model.weights.best.hdf5
Epoch 9/20
1505/1505 [=====] - 7s 5ms/step - loss: 2.0681 - acc: 0.3203 - val_loss: 2.3915 - val_acc: 0.2300

Epoch 00009: val_loss improved from 2.39367 to 2.39149, saving model to model.weights.best.hdf5
Epoch 10/20
1505/1505 [=====] - 7s 5ms/step - loss: 1.8587 - acc: 0.3967 - val_loss: 2.3957 - val_acc: 0.2400

Epoch 00010: val_loss did not improve from 2.39149
Epoch 11/20
1505/1505 [=====] - 7s 5ms/step - loss: 1.6446 - acc: 0.4950 - val_loss: 2.3905 - val_acc: 0.2500

Epoch 00011: val_loss improved from 2.39149 to 2.39050, saving model to model.weights.best.hdf5
Epoch 12/20
1505/1505 [=====] - 7s 5ms/step - loss: 1.4765 - acc: 0
```

```
In [19]: # load the weights that yielded the best validation accuracy
model.load_weights('model.weights.best.hdf5')
```

```
In [20]: # evaluate and print test accuracy
score = model.evaluate(X_test, y_test, verbose=0)
print('\n', 'Test accuracy:', score[1])
```

Test accuracy: 0.24378109474976858